

ANUVA NEGI

anuvanegi@gmail.com • (469) 432-1606 • Arlington, TX • [linkedin.com/in/anuva-negi](https://www.linkedin.com/in/anuva-negi) • github.com/silverfrost702 • [Portfolio](#)

PROFESSIONAL SUMMARY

Data Science graduate student and Dean's Scholar with 2+ years of experience in data engineering and analytics. Proficient in Python, SQL and AI-assisted tools, building ETL pipelines, predictive models and interactive dashboards. Delivers actionable insights that drive strategic decision-making and cost savings for content and product teams.

SKILLS

Programming Languages: Python, SQL, R
Machine Learning & Analytics: Predictive Modeling, Time-Series Forecasting, EDA, Data Analysis, Feature Engineering, Fraud Detection
Libraries & Frameworks: Pandas, NumPy, Scikit-learn, TensorFlow, CatBoost, XGBoost, Statsmodels, Flask, PyTorch, Keras
Data Visualization & AI: Power BI, Tableau, Plotly, Streamlit, Statsmodels, Large Language Models (LLMs), Generative AI, RAG
Databases & Tools: Snowflake, MySQL, MongoDB, Airflow, Informatica, Git, JIRA, Confluence, AWS S3, PostgreSQL
Certifications: SnowPro Core Certification, Google Data Analytics Professional Certificate

EDUCATION

M.S. in Data Science | **University of Texas at Arlington** | GPA: 4.0 / 4.0

Aug 2024 - May 2026 | USA

B.Tech in Computer Science & Engineering | **Birla Institute of Applied Sciences**

Aug 2018 - Jun 2022 | India

EXPERIENCE

Data Engineer | **Genpact (PIMCO)**

May 2023 - Jul 2024

- Spearheaded migration of 5TB+ datasets from Oracle and MySQL to Snowflake, contributing to \$5M+ annual infrastructure cost savings.
- Designed and executed 100+ ETL pipelines using Python, SQL, Perl, Informatica, Airflow, and Autosys with 100% on-time delivery.
- Boosted query performance by 30% through production code development and model redesign, reducing latency for downstream teams.
- Collaborated with data architects and business stakeholders to document data lineage, schema definitions, and pipeline dependencies, improving team onboarding and audit readiness and enabling feature engineering and predictive modeling.

Business Analyst | **Genpact (PIMCO)**

Sep 2022 - Apr 2023

- Produced 50+ analytical reports and interactive dashboards using Power BI and SQL used directly in executive leadership reviews, influencing strategic planning decisions.
- Bridged gap between quant and dev teams by converting 150+ business requirements into structured JIRA stories, reducing rework.
- Conducted gap analysis across workflows and business objectives, identifying inefficiencies and delivering Python-based dashboards and exploratory data analysis solutions adopted by 2 delivery teams.

Software Engineer Intern | **Enquero**

Feb 2022 - Aug 2022

- Completed 50+ hours of Teradata and cloud data platform training; conducted EDA using Pandas, Matplotlib and Seaborn.
- Built 3 internal web-based interfaces using HTML, CSS, and JavaScript, enabling non-technical stakeholders to interact with EDA reports.

PROJECTS

ChurnSight: Customer Churn Prediction & Retention Analytics | XGBoost, SHAP, Plotly, Imbalanced-learn

[GitHub](#) | [Demo](#)

- Engineered an end-to-end churn prediction pipeline on the IBM Telco dataset (7K+ customers) using Python for data extraction and preprocessing, creating 6 SHAP-validated features and training an XGBoost model achieving ROC-AUC 0.847 and Recall 0.840 (CV AUC 0.845 $\pm$ 0.011).
- Optimized retention targeting using business loss metrics (FN=\$1,500, FP=\$200), reducing intervention costs by \$14.9K and projecting \$740K+ annual savings across a 70K-customer portfolio.
- Delivered an interactive Streamlit dashboard built with Plotly that provides single-customer SHAP waterfall explanations, batch risk-tier segmentation, and an audit trail for modeling decisions, automating narrative summaries with LLM prompts.

SmogCast – Air Quality Classification & Forecasting | CatBoost, LSTM, Transformer, Streamlit

[GitHub](#) | [Demo](#)

- Designed end-to-end ML pipeline on 700K+ hourly sensor records across 26 Indian cities; achieved 93% multiclass AQI classification accuracy using CatBoost with SMOTE for class imbalance and lag/rolling-window feature engineering on time-series data.
- Implemented a 7-day PM2.5/PM10 forecasting system using Bi-LSTM and Transformers (9–12 RMSE) and integrated SHAP explainability, deploying a Streamlit dashboard with WHO alerts and policy recommendations.
- Developed an interactive dashboard using Plotly and D3.js to visualise hourly AQI classifications and 7-day PM2.5/PM10 forecasts.

Predictive Analytics for Academic Incident Management | Pandas, Statsmodels, Streamlit, Plotly

[GitHub](#) | [Demo](#)

- Extracted and normalized ServiceNow incident logs, engineering an ITSM analytics pipeline and applying temporal aggregation, anomaly detection and seasonality analysis to forecast workload spikes.
- Created interactive Streamlit dashboard identifying service hotspots and infrastructure strain, enabling IT leadership to make data-driven staffing decisions and plan targeted hardware refreshes.
- Implemented data quality checks and anomaly filtering routines to address inconsistencies, improving pipeline reliability and accuracy

Pokémon Data Showdown | Python, Scikit-learn, Streamlit

[GitHub](#) | [Demo](#)

- Placed 2nd out of 20 teams in the competition, recognized for model creativity and deployment quality.
- Applied clustering, feature scaling, dimensionality reduction, and silhouette scoring to identify role-based archetypes and optimize teams.
- Published an interactive Streamlit app enabling exploratory stat analysis and team-building recommendations.

ACHIEVEMENTS

Student Employee of the Year, Award of Excellence, University of Texas at Arlington

Mar 2026

Graduate Research Assistantship, Merit-based 1-year funding award, UTA

Aug 2025